A nonlinear MPC approach to minimize toxicity in HIV-1 infection multi-drug therapy [*]

João M. Lemos [*] **João V. Pinheiro** [**] **Susana Vinga** [***]

[*] *INESC-ID/IST/Univ. Técnica de Lisboa, Portugal (jlml@inesc-id.pt).*
[**] *INESC-ID, Portugal (jvpinheiro@gmail.com).*
[***] *INESC-ID, Portugalsvinga@kdbio.inesc-id.pt*

Abstract: This work addresses the problem of minimizing the toxicity of the therapy applied to control HIV-1 infection when a multi-drug therapy is (as usually) employed. An approach based on nonlinear model predictive control (MPC), in which the input is a train of impulses corresponding to periodic pill administration to the patient, is followed. Patient adherence to treatment and their impact on virus drug resistance is considered in the model used. The control objective consists of driving the viral load to a low specified value while minimizing the amount of drugs administered to the patient. Since various dose combinations of reverse transcriptase inhibitor (RTI) and protease inhibitor (PI) drugs can be given depending on the weights of the cost function minimized by periodic MPC while attaining the same control objective, it is proposed that these weights can be adjusted to minimize the toxicity of the drug cocktail administered. This approach is illustrated by means of simulation in a model that incorporates both pharmaco-kinetics and pharmaco-dynamics of the drugs considered. *Copyright CONTROLO 2012.*

Keywords: Nonlinear Model Predictive Control, HIV-1 infection control, immunology, toxicology, biomedical systems.

1. INTRODUCTION

1.1 Motivation

In addition to their desired therapeutical effects, all drugs have a toxic effect that is enhanced above a given threshold. When adjusting the drug doses, the values selected should be such that they are large enough to create a therapeutic effect, while at the same time minimizing the toxic effect. In the therapy of HIV-1 infection a combination of drugs is administered to the patient in order to reduce the toxic effects. This problem is complicated by the fact that the human organism is a dynamical system in the sense that the response to a set of drugs depends not just on instantaneous dose values but on the state of the patient as well as on the past treatment history. Current work on modeling of the interaction between drugs and HIV-1 infection motivate therefore that the problem of finding the therapy with the smallest toxic effect is formulated as a dynamic optimization problem. Model predictive control provides a powerful tool to address this problem.

1.2 Literature review

A recent paper addresses the application of periodic control to drug delivery. Since the seminal paper Perelson and Nelson (1999) that pointed the attention to the importance of phenomena that take place in time scales of days and weeks, that are fast when compared with the time scale of AIDS, many works have been reported in the literature concerning both model development and analysis and therapy design using control algorithms Chang and A. Astolfi (2009a,b). Other examples

of control techniques used include nonlinear control Ge *et al.* (2005); Barão and Lemos (2007); Mhawej *et al.* (2010) and optimal control Kutch and Gurfil (2002); Stengel (2008).

Predictive control is currently receiving an increased attention in relation to HIV-1. In Zurakowski and Teel (2006) MPC is used to schedule interruptions in highly active anti-retroviral therapy (HAART) used to simulate a therapeutic vaccine. Treatment schedules based on robust multirate MPC are proposed in Elaiw and Xia (2009). In order to minimize drug consumption, Pannocchia *et al.* (2010) proposes a MPC based algorithm in which the dose (given with a sampling time of one week) is restricted either to be zero or the maximum acceptable value. In Wu *et al.* (2006) a model of pharmacodynamics of antiretroviral drugs in HIV-1 infected patients that incorporates drug susceptibility and patient adherence is discussed.

1.3 Paper contributions and structure

This work addresses the problem of minimizing the toxicity of the therapy applied to control HIV-1 infection when a multi-drug therapy is employed. An approach based on nonlinear model predictive control (MPC) is proposed in which the weights of the cost function are adjusted.

The paper is organized as follows: After the Introduction (this section) that motivates the problem, performs a concise literature review and describes the paper contributions and structure, the HIV-1 infection model used is presented in section 2. This model is exploited in section 3 to address nonlinear state estimation, required for the control algorithm used, and in section 4 to develop a nonlinear model predictive control (NMPC) algorithm with periodic inputs. Designing the controller for minimum toxicity is addressed in section 4. Finally, section 6 draws conclusions.
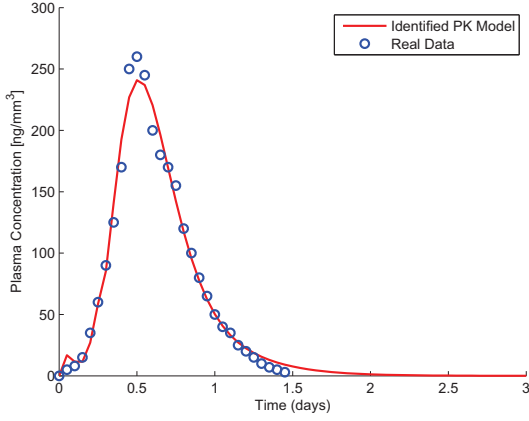
Fig. 2. RTI pharmacokinetics model impulse response.

2. PK/PD MODEL FOR HIV-1 INFECTION

Figure 1 shows the model of HIV-1 infection and drug interaction considered in this work, including a model of patient adherence to the therapeutics, drug pharmacokinetics and pharmacodynamic model. The pharmacodynamic model includes a model of the development of virus resistance to drugs, for patients without perfect adherence that is a variant of the one presented in Wu *et al.* (2006). The pharmacodynamic model includes a model of the development of resistances in response to non-perfect adherence.

It is assumed that two different types of drugs are being administered to the patient: A reverse transcriptase inhibitor (RTI, drug number 1) and a protease inhibitor (PI, drug number 2). Hence, corresponding to each drug, there is one pharmacokinetics and one pharmacodynamic model, corresponding to indexes 1 and 2 in the block diagram of figure 1.

The inputs to the the model are the drug doses of each type, u_{PK1} for the RTI dose and u_{PK2} for the PI dose and the patient adherence A. The drug doses are assumed to be trains of square pulses of very small duration and with an height proportional to the dose taken. The adherence A is defined as the probability that the patient will actually take the dose prescribed by the controller. Hence, at each discrete time k, the adherence model generates a random number α_A with a uniform distribution between 0 and 100. If $\alpha_A \leq A$ then the dose effectively taken of drug i, at time k, $\bar{u}_{PKi}(k)$ is made equal to $u_{PKi}(k)$; otherwise, it is made equal to 0.

The pharmacokinetics (PK) model of drug i (block PK_i in figure 1) is a linear, time invariant, model, with 3 poles, 3 zeros and unit delay. These PK models relate the amounts of drug administered to the patient with the drug plasma concentration c_{p1} and c_{p2}. Figure 2 shows the impulse response of the PK model used for the RTI drug. The impulse response used for PI is similar. These are typical curves, adapted from published drug PK data with modifications and taken for exemplificative purposes. Of course the results depend on the specific drugs considered.

The pharmacodynamic (PD) model consists of two cascade parts (figure 1). The first part relates the plasma concentration of each drug with the corresponding drug effect (u_1 for RTI and u_2 for PI). This is assumed to be the static nonlinear relation known as the "Hill equation" that is given by
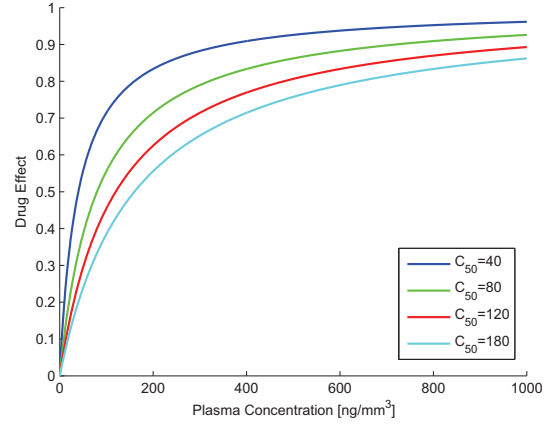

Fig. 3. The relation between the plasma concentration and the drug effect as given by the Hill equation.
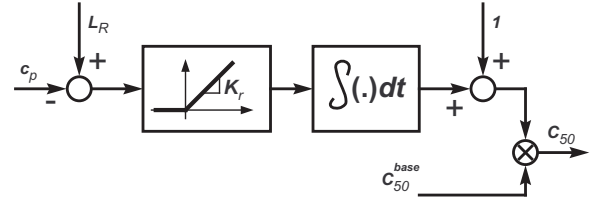

Fig. 4. Resistance model of HIV-1.

$$u_i = \frac{c_{pi}}{c_{50}^i + c_{pi}}, \tag{1}$$

where c_{50}^i is a parameter for each i. For drug i, when the plasma concentration is $c_{pi} = 0$ the corresponding effect u_i is zero. When the plasma concentration c_{pi} grows, the effect u_i approaches 1. For $c_{pi} = c_{50}^i$, the effect is $u_i = 0.5$.

Figure 3 shows the relation between the plasma concentration and the drug effect as given by the Hill equation for several values of the parameter C_{50}. There is one of these curves for each drug considered. When C_{50} grows the patient becomes more insensitive to the drug, meaning that an higher plasma concentration is needed to achieve the same effect.

The resistance model tries to represent the effect that low plasma drug concentrations induce virus mutations that render the virus more resistant to the drug used. The increase of virus resistance means that the same effect can only be achieved with an higher plasma concentration. A way of representing this effect consists in increasing the c_{50}^i parameter depending on the periods in which the plasma concentration is below the threshold L_R that allows virus mutation Wu *et al.* (2006). This is the resistance model that is detailed in figure 4 (for simplicity the index of the drug has been dropped in this figure).

It is assumed that the parameter C_{50} increases proportionally to the area below L_R and limited by the plasma concentration c_p. Whenever the plasma concentration decreases below L_R this area will increase, causing the factor f and, consequently, C_{50} to increase. In compact terms, this reads

$$C_{50}(t) = f(t)C_{50}^{base} \tag{2}$$

where C_{50}^{base} is a constant giving the initial value of C_{50} and $f(t)$ is given by
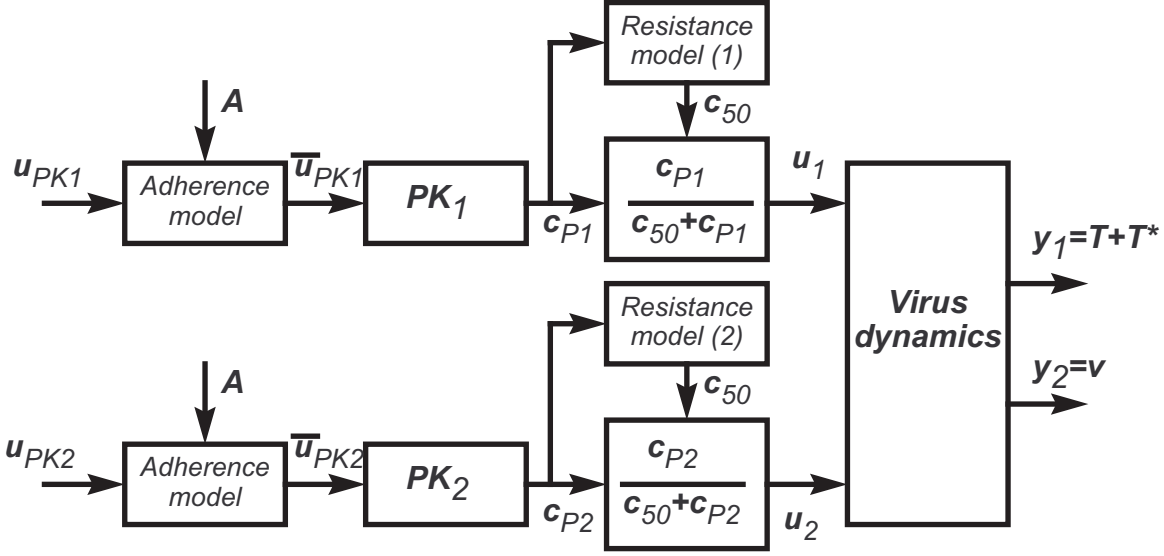
Fig. 1. Pharmacokinetic/pharmacodynamic model of HIV-1 infection incorporating adherence and resistance models.

$$f(t) = 1 + K_r \int_0^t \max[0, L_r - c_p(\tau)]d\tau \qquad (3)$$

where K_r is a parameter that measures the capacity of the virus to develop resistance in the presence of a drop of drug plasma concentration. For simplicity of notation the index of the drug has been omitted.

The second block of the PD model represents the interaction between drug effects u_1 and u_2 and virus dynamics. A simple model is used that considers only the following three state variables Perelson and Nelson (1999):

- x_1, plasma concentration of healthy T-CD4+ cells;
- x_2, plasma concentration of infected T-CD4+ cells;
- x_3, plasma concentration of free virus particles (virions).

A balance, together with simple assumptions on interaction rates between the different elements involved yields the following nonlinear state space model:

$$\begin{cases} \dot{x}_1 = s - dx_1 - (1-u_1)\beta x_1 x_3 \\ \dot{x}_2 = (1-u_1)\beta x_1 x_3 - \mu x_2 \\ \dot{x}_3 = (1-u_2)kx_2 - cx_3 \end{cases}$$
$$y = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \qquad (4)$$

where β, μ, k and c are patient dependent parameters.

For the sake of simplicity in mathematical manipulations, model (4) can be written in the following condensed form

$$\begin{aligned} \dot{x} &= \phi(x(t), u(t)) \\ y &= h(x, u) \end{aligned} \qquad (5)$$

with obvious definitions for functions ϕ and h.

Since the control algorithms considered operate in discrete time, the above model is approximated using the Euler method. Let Δ be the sampling interval. With an obvious definition for Φ, the discrete model is written as

$$x(k+1) = \Phi(x(k), u(k)) \qquad (6)$$

3. NONLINEAR STATE ESTIMATION

The NMPC algorithm used requires knowledge of the state variables. Since not all of these are directly measured, the Extended Kalman Filter (EKF) is used to produce state estimates according to the following equations:

For prediction:

$$\hat{x}_{k|k-1} = \Phi(\hat{x}_{k-1|k-1}, u_{k-1}) \qquad (7)$$
$$P_{k|k-1} = F_k P_{k-1|k-1} F_k^T + Q_k \qquad (8)$$

with the jacobian given by

$$F_{k-1} = \frac{\partial \Phi}{\partial x}|_{\hat{x}_{k-1|k-1}, u_{k-1}} =$$

$$= \begin{bmatrix} 1 - \Delta(d + U_1)\beta x_3 & 0 & -\Delta U_1 \beta x_1 \\ \Delta U_1 \beta x_3 & 1 - \Delta\mu & \Delta U_1 \beta x_1 \\ 0 & \Delta U_2 k & 1 - \Delta c \end{bmatrix} \qquad (9)$$

where $U_1 = 1 - u_1$, $U_2 = 1 - u_2$ and Δ is the sampling interval chosen to be a submultiple of the one used for control.

For update:

$$\tilde{y}_k = z_k - h(\hat{x}_{k|k-1}) \qquad (10)$$
$$S_k = H_k P_{k|k-1} H_k^T + R_k \qquad (11)$$

with

$$H_{k-1} = \frac{\delta h}{\delta x}|_{\hat{x}_{k-1|k-1}} = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \qquad (12)$$

and

$$\begin{aligned} K_k &= P_{k|k-1} H_k^T S_k^{-1} \\ \hat{x}_{k|k} &= \hat{x}_{k|k-1} + K_k \tilde{y}_k \end{aligned} \qquad (13)$$
$$P_{k|k} = (I - K_k H_k) P_{k|k-1} \qquad (14)$$

The noise variances are selected as

$$Q_k = \begin{bmatrix} 0.5 & 0 & 0 \\ 0 & 0.001 & 0 \\ 0 & 0 & 50 \end{bmatrix} \quad R_k = \begin{bmatrix} 0.5 & 0 \\ 0 & 0.0005 \end{bmatrix} \qquad (15)$$

and the initial error covariance matrix is $P_{0|0} = I$. These covariance values were taken as "tuning knobs", being adjusted so as to optimize the estimates.

A multi-rate version of EKF is used. The prediction stage is used in every iteration while filtering is applied only in instants that are multiple of this time unit.

4. PERIODIC NONLINEAR MPC OF HIV-1 INFECTION

The control problem consists in bringing down the viral load, that corresponds to the state x_3, to a level of 50 $copies/mm^3$ and keeping it there, in less than 56 days. The manipulated variables are the drug doses u_{PK1} and u_{PK2}

The NMPC control is obtained by solving in a receding horizon way, at each current discrete time t, the following constrained minimization problem:

$$\min_{\bar{u}_t^{t+T_c-1}} J(x(t), \bar{u}_t^{T_c}, T_c, T_p) \tag{16}$$

where T_c is the control horizon and T_p is the prediction horizon, verifying $T_p \geq T_c$. The sequence of virtual control $\bar{u}$ is

$$\bar{u}_t^{t+T_c-1} := [\bar{u}(t) \dots \bar{u}(t+T_c-1)]^T$$

The receding horizon cost J is defined as

$$J(x(t), \bar{u}_t^{T_c}, T_c, T_p) = \sum_{k=1}^{T_p} L[\bar{x}(t+k), \bar{u}(t+k-1)] \tag{17}$$

where the stage cost is defined by

$$L(x, u) = (x-x_s)^T Q(x-x_s) + u^T Ru$$

with

$$Q = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} R = \begin{bmatrix} w_{u_1} & 0 \\ 0 & w_{u_2} \end{bmatrix} \tag{18}$$

and x_s is the reference, given by

$$x_{s,3}(t) = 50 + (902.8 - 50)e^{-\tau_r t} \tag{19}$$

The variable τ_r is the time constant that determines the speed at which the reference signal decreases towards the desired value of 50 $copies/mm^3$, and was selected as $\tau_r = 0.2\ days^{-1}$.

It is assumed that, for $i = 1, \dots, T_p - T_c$

$$\bar{u}(T_c + i) = \bar{u}(T_c) \tag{20}$$

The virtual state trajectory satisfies the system dynamics with initial condition $\bar{x}(t)$ given by the actual state x at time t. Since the actual state is not available for direct measurement, it is replaced by its EKF estimate, $\hat{x}$. Thus, for $k = 0, \dots, T_{p-1}$:

$$\bar{x}(t+k+1) = \Phi(\bar{x}(t+k), \bar{u}(t+k)) \tag{21}$$

starting from

$$\bar{x}(t) = \hat{x}(t) \tag{22}$$

According to a receding horizon strategy, once the problem (16) is solved with respect to the virtual control $\bar{u}$, the value of the manipulated variable to actually apply to the plant at time t is $u(t) = \bar{u}(t)$. The same procedure is again repeated at time $t+1$. Figures 5 and 6 show the viral load and the drug dosage for a patient with an adherence of 95% controlled with NMPC. Figure 7 shows the corresponding plasma concentration of RTI and PI. As seen in figure 6 the doses of both drugs applied grow initially such as to drive the viral load to a low level. Then, the doses are reduced, to keep the viral load where desired. After the fail of a control sample, the controller raises the value of subsequent samples to compensate the drop of drug concentration in plasma.

Figure 8 shows the dependence of the cost on the prediction horizon. It should be remarked that the horizon actually "seen"
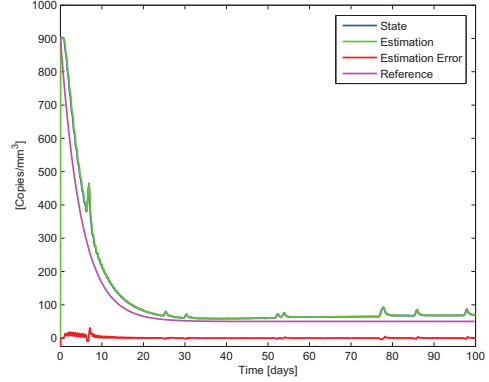

Fig. 5. Viral load response for a patient with adherence of 95% and $K_r = 1 \times 10^{-4}$.
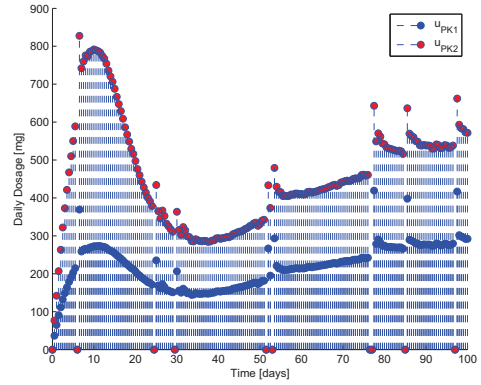

Fig. 6. Dosage for a patient with adherence of 95% and $K_r = 1 \times 10^{-4}$.

Fig. 7. Drug plasma concentration in the situation of figures 5 and 6.

by the algorithm is the product of T_p by the sampling interval. Smaller sampling intervals require a bigger T_p so that the same performance is achieved. For a given sampling rate, when T_p increases, the performance increases. However, this increase is smaller and smaller as T_p gets bigger. Due to plant model mismatches, it is expected that, when T_p grows, after the initial decrease of the cost this will increase again, for larger values of
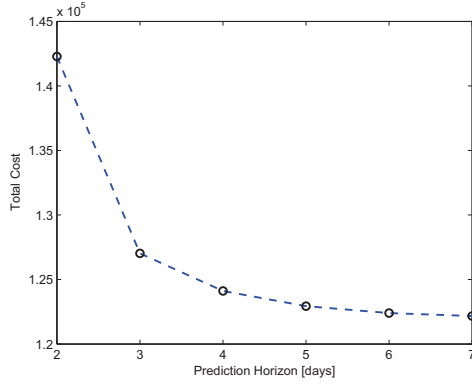
Fig. 8. Cost and Average Error Values for Varying Prediction Horizons

T_p, due to the degradation of prediction performance over long time horizons.

5. COST WEIGHTS AND TOXICITY

Toxicity is optimized by adjusting the weights w_{u_1} and w_{u_2} on the manipulated variables (doses of RTI and PI) in the cost function (17). This is done as explained below.

The plots in figure 9 show the effect of the cost weight parameter w_{u_1}. Plots for the influence of w_{u_2} are similar. When w_{u_1} increases there is a decrease in the average of the amount of RTI that is administered to the patient and an increase in PI. Thus, the parameter w_{u_1} regulates a trade-off between these two types of drug (shown in figure 10) that can be explored to minimize toxicity of the overall treatment.

Denote by $< u_{PKi} >$ the average value of u_{PKi} and by g_i the function that measures the toxicity for drug i, $i = 1, 2$. Figures 11 and 12 represent graphically the functions g_i for RTI and PI. These functions yield a relatively small value for the toxicity up to a threshold value of the drug dose and then grow very fast.

Using the above model for the toxicity of each drug, the combined average toxicity of both dugs is given by

$$T_x = g_1(< u_{PK1} >) + g_2(< u_{PK2} >).$$

On the other way, using NMPC, the average values of both drugs are related by the function plotted in figure 10, that we call Υ. This function is found by eliminating w_{u_1} between both curves in figure 9. We have thus

$$< u_{PK2} >= \Upsilon(< u_{PK1} >) \tag{23}$$

and the total toxicity due to both drugs is given by

$$T_x = g_1(< u_{PK1} >) + g_2(\Upsilon(< u_{PK1} >)), \tag{24}$$

a function that has the shape depicted in figure 13.

Hence, we can find a value for the weight w_{u_1} that minimizes treatment toxicity, by finding the function u_{PK1} that solves the optimization problem

$$\min_{< u_{PK1} >} \{g_1(< u_{PK1} >) + g_2(\Upsilon(< u_{PK1} >))\} \tag{25}$$

and then computing the corresponding value for w_{u_1} from figure 10.

6. CONCLUSIONS AND DISCUSSION

The use of a periodic nonlinear MPC, together with a model incorporating drug PK and PD, virus dynamics and patient

Fig. 9. Average input Values as a function of the input weight w_{u_1}.

Fig. 10. Selecting the average input values for u_{PK1} and u_{PK2} by varying the value of the input weight w_{u_1}
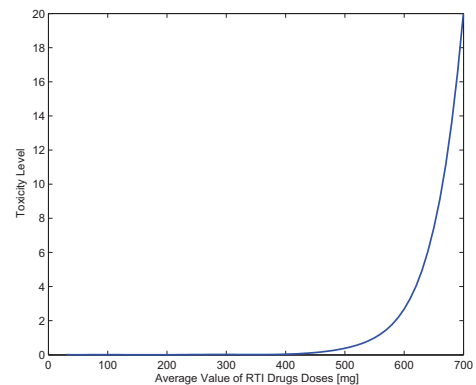

Fig. 11. Toxicity of RTI as a function of average drug dose.

adherence provides an adequate way to compute optimal drug dosage for HIV-1 infection therapy. For a patient adherence of 90% the controller could still meet the specifications. The cost dependence on controller parameters as well as on model parameters have been studied, including parameter mismatching. The choice of weights in the cost function minimized by periodic NMPC provides an extra degree of freedom that is exploited to optimize therapy toxicity while attaining the control objectives.

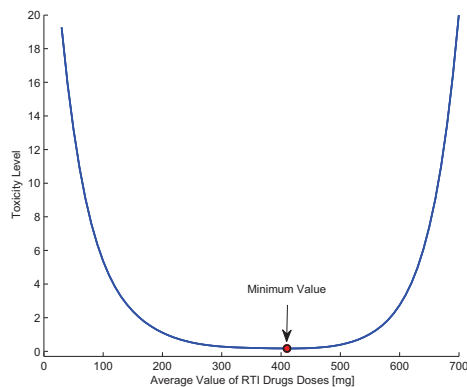Fig. 12. Toxicity of PI as a function of average drug dose.

Fig. 13. Combined toxicity of both RTI and PI as a function of average drug dose.

REFERENCES

Barão, M. and J.M. Lemos (2007). Nonlinear control of HIV-1 infection with a singular perturbation model, *Biomedical Signal Processing and Control* 2:248-257.

Chang, H. and A. Astolfi (2009). Control of HIV infection dynamics: Approximating high-order dynamics by adapting reduced order model parameters. *IEEE Control Systems Mag.*, 28(April):28-39.

Chang, H. and A. Astolfi (2009). Enhancement of the immune system in HIV dynamics by output feedback. *Automatica* 45:1765-1770.

Elaiw, A. M. and X. Xia (2009). HIV dynamics: Analysis and robust multirate MPC-based treatment schedules, 359:285-301.

Ge, S. S.; Z. Tian and T. H. Lee (2005). Nonlinear control of a dynamic model of HIV-1. *IEEE Trans. Biomed. Eng.*, 52(3):353-361.

KS05 Knorr, A. L. and R. Srivastava (2005). Evaluation of HIV-1 kinetic models using quantitative discrimination analysis. *Bioinformatics*, 21(8): 1668-1677.

Kutch, J. J. and P. Gurfil (2002). Optimal control of HIV infection with a continuously-mutating viral population. *Proc. Amer. Control Conf.*, Anchorage, AK May 8-10, 2002. 4033-4038.

Mhawej, S. S.; C. H. Moog, F. Biafore and C. Brunet-Franois (2010). Control of the HIV infection and drug dosage. *Biomedical signal processing and control*, 5:45-52.

Pannocchia, G.; M. Laurino and A. Landi (2010). A model predictive control strategy toward optimal structured treatment interruptios in anti-HIV therapy. *IEEE Trans. Biomed. Eng.*, 57(5):1040-1050.

Perelson, A. S. and P. W. Nelson (1999). Mathematical analysis of HIV-1 dynamics in vivo. *SIAM Review*, 41(1):3-44.

Stengel, R. F. (2008). Mutation and control of the human immunodeficiency virus. *Mathematical Biosciences*, 213:93-102.

Varigonda, S.; T. T. Georgiou, R. A. Siegel and P. Daoutidis (2008). Optimal periodic control of a drug delivery system. *Computers and Chemical Engineering*, 32:2256-2262.

Wu, H.; Y. Huang, E. P. Acosta *et al.* (2006). Pharmacodynamics of antiretroviral agents in HIV-1 infected patients: using viral dynamic models that incorporate drug susceptibility and adherence. *J. of Pharmacokinetics and Pharmacodynamics*, 33(4):399-413.

Zurakowski, R. and A. R. Teel (2006). A model predictive control based scheduling method for HIV therapy. *J. of Theoretical Biology*, 238:368-382.